CODEX + CLAUDE / WINDOWS GATEWAY

One desktop. Two model providers.

OpenAI stays configured as the default; Claude requests use your Anthropic API credits.

Use Work > On your computer. Cloud-mode chats use the app's separate OpenAI catalog.

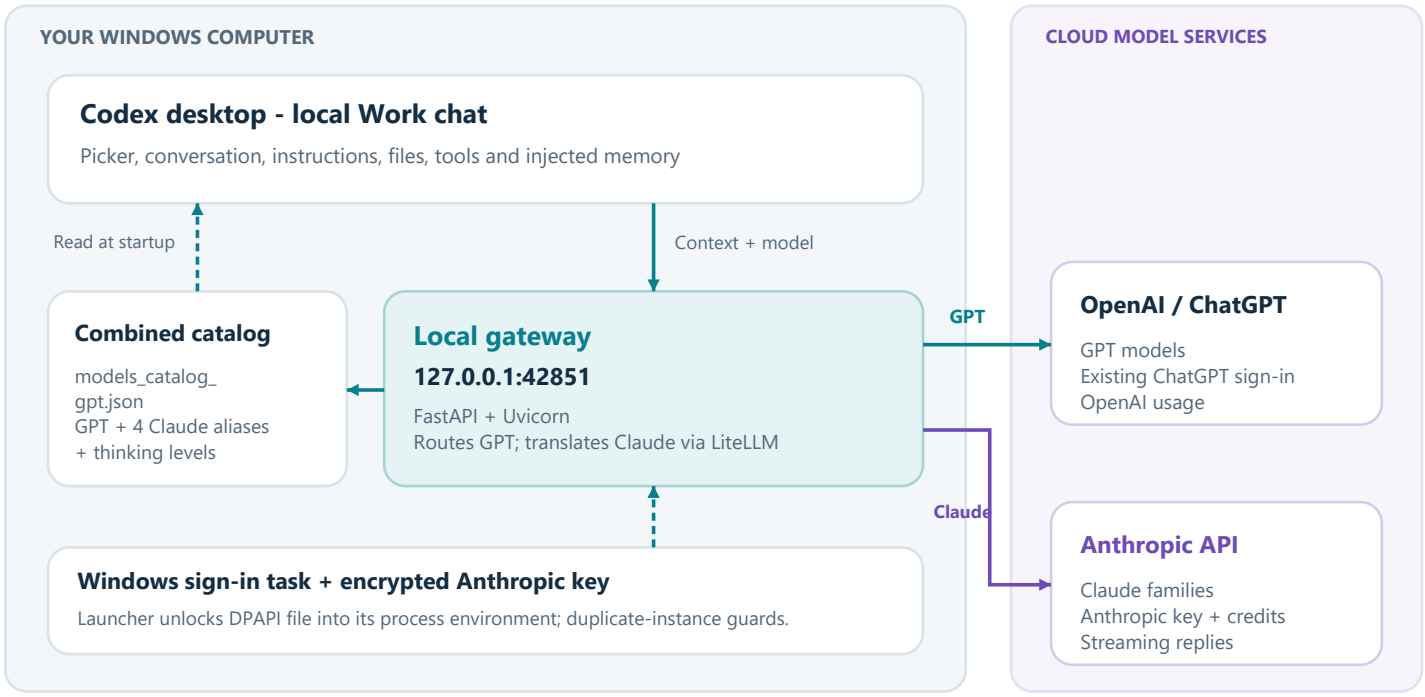

All models run in the cloud. "On your computer" refers to the agent, tools and local state.

HOW THE REQUEST WORKS

1 Assemble context

Codex builds the prompt from chat history or a summary, instructions, selected files, tool results and any injected local memories.

2 Route and stream

GPT keeps ChatGPT authentication. Claude uses only the Anthropic key. Aliases pin versions per chat; LiteLLM translates requests and streams replies.

3 Refresh model choices

Claude models refresh hourly. GPT catalogs refresh on demand. New chats use the newest Claude version; new families or picker capabilities need a Codex restart.

CODE AND CONFIGURATION

config.toml	GPT default + gateway URL + catalog path; request compression disabled.
gateway_gpt.py	Routing, model refresh, version pins, context forwarding and response streams.
models_catalog_gpt.json	Explicit GPT + Claude picker catalog; ordered with GPT first.
start_claude_picker_gpt.ps1	Unlocks the saved key and starts the loopback server quietly.
compaction_gpt.py	Windows-encrypted Claude summaries for long-chat continuation.
install_autostart_gpt.ps1	Configures the user-logon task; restart policy and hidden execution.

Source: your stable project checkout; runtime data stays outside Git

Data: %LOCALAPPDATA%\CodexClaudeGateway (key, catalogs, pins, venv)

BOUNDARIES TO REMEMBER

The gateway must run for local GPT and Claude requests. Key expiry still applies. Any context included in a Claude request goes to Anthropic. The app remembers new-composer selections. GPT-encrypted compaction cannot be decoded by Claude, so switching a long GPT chat can fail.